
Beyond Absolute Imitation: Anchored Residual Guidance for Privileged On-Policy Distillation

Anonymous Author(s)

Affiliation

Address

email

Abstract

On-policy distillation (OPD) has demonstrated strong empirical gains in enhancing complex reasoning in LLMs by aligning a student model with a teacher’s predictive distribution over the student’s own trajectories. An emerging variant, Privileged OPD, further strengthens this paradigm by employing a self-teacher model augmented with privileged information, such as oracle traces, to mitigate teacher-student capacity gaps while providing dense, answer-directed supervision. However, current methods treat privileged information as a monolithic imitation target, failing to disentangle locally reachable reasoning steps from future-conditioned oracle signals. Consequently, the student is encouraged to match a hindsight-biased distribution that often falls outside its local predictive support. This reachability mismatch incentivizes the student model to skip valid intermediate reasoning in favor of locally unsupported shortcuts. To resolve this, we introduce **Anchored Residual On-Policy Distillation (AR-OPD)**, a dual-view framework that disentangles privileged supervision. Rather than enforcing strict full-view imitation, AR-OPD establishes a locally compatible anchor using a partially privileged teacher, isolating and injecting oracle foresight as a controlled residual to provide destination-directed guidance. Across diverse reasoning tasks, **AR-OPD outperforms full privileged OPD by 2.3 points and SFT by 7.9 points**. Crucially, this anchored residual mechanism reduces hindsight leakage by up to **23%** and mitigates late-stage drift, yielding up to a **7.2-point advantage** on challenging long-horizon trajectories exceeding 768 tokens.

1 Introduction

On-policy distillation (OPD) has rapidly emerged as an effective paradigm for large language model post-training, driving substantial gains in recent industry pipelines [Qwen Team, 2025, Xiaomi MiMo Team, 2025, GLM-5 Team, 2026]. By querying a teacher on states actively induced by the student, OPD addresses the off-policy exposure bias of supervised fine-tuning [Brown et al., 2020, Wei et al., 2022, Ouyang et al., 2022, Taori et al., 2023] and the reward sparsity of reinforcement learning [Schulman et al., 2017, Guo et al., 2025, Dong et al., 2024, Sheng et al., 2024, Thinking Machines Lab, 2025, Agarwal et al., 2024]. However, standard cross-model OPD often suffers from pattern mismatch: distilling a structurally incompatible teacher can degrade student performance [Li et al., 2026, Fu et al., 2026, Jang et al., 2026].

To circumvent this, *Privileged OPD* (e.g., OPSD, SDFT [Zhao et al., 2026]) restricts the teacher and student to the same base model but augments the teacher with auxiliary context, such as oracle traces. By allowing the self-teacher to “see more,” privileged OPD improves teacher-student compatibility and provides dense supervision aligned with the student’s native generation manifold. While Privileged OPD densifies the learning signal, it can also shift the teacher’s distribution away from the student’s causal predictive support.




Figure 1: Privileged-information leakage as an off-support target. A representative example illustrating privileged-information leakage in a privileged OPD-trained model, where the model appeals to an invisible reference solution during inference. Because the full privileged teacher conditions on oracle information unavailable to the student, it can assign probability mass to answer-conditioned states. Direct imitation therefore turns factually correct privileged information into a shortcut target.

38 This creates a different failure mode from ordinary teacher-student mismatch. The teacher is not
 39 merely stronger or weaker; it is conditioned on information that the student cannot access at the same
 40 prefix. Thus, the central question is not whether privileged information is useful, but how much of it
 41 remains locally learnable as a token-level target.

42 Consequently, treating this fully privileged teacher as an absolute imitation target introduces a subtle
 43 but persistent optimization pathology. Because the full-view oracle conditions on the complete future
 44 context, its distribution implicitly incorporates hindsight-derived information and future-conditioned
 45 structure. It shifts probability mass toward tokens that, while factually correct, remain causally
 46 unsupported from the student’s current prefix—such as anticipating a final answer branch before
 47 establishing the necessary intermediate rationale. We characterize these premature emissions as
 48 **privileged-information leakage** or *shortcut* events (Figure 1 illustrates this target-side failure, where
 49 a privileged OPD-trained model appeals to an invisible “reference solution” during inference, a
 50 phenomenon later quantified in Figure 6). In Section 4, we show that this mismatch becomes most
 51 visible near the rollout tail, where full-view targets assign increasing mass outside the student’s local
 52 support. We further show that full-view reliability itself degrades with privileged-context length,
 53 making monolithic full-view imitation brittle precisely in long-horizon reasoning.

54 To overcome this absolute imitation trap, we propose **Anchored Residual On-Policy Distillation**
 55 (**AR-OPD**), an objective designed to extract privileged guidance while preserving target reachability.
 56 Rather than forcing the assimilation of a monolithic full teacher, AR-OPD disentangles privileged
 57 supervision into two controllable components: a locally compatible **ANCHOR** and a destination-
 58 directed **RESIDUAL**. The **ANCHOR** uses a partially privileged teacher to provide a causally reachable
 59 target that isolates the student from hindsight-biased answers, while the **RESIDUAL** injects the full-
 60 view teacher’s marginal foresight as a λ -scaled log-probability update over the anchor. This shifts
 61 imitation from an unsupported hindsight coordinate to a controlled, future-directed update.

62 Our contributions are summarized as follows:

- 63 ① **Controlling the impact of privileged information.** We show that monolithic use of oracle-
 64 augmented teachers can induce optimization pathologies, and introduce a framework that extracts
 65 useful destination-directed foresight while controlling hindsight-biased manifold drift.
- 66 ② **Anchored Residual On-Policy Distillation (AR-OPD).** We introduce a dual-view distillation
 67 method that operationalizes this control through a causally reachable partial-oracle **ANCHOR**
 68 and a λ -scaled log-space **RESIDUAL**, transferring privileged guidance without sacrificing local
 69 sequence compatibility.
- 70 ③ **Mechanistic evidence and cross-domain gains.** AR-OPD suppresses privileged-information
 71 leakage, reducing shortcut events by **21.7%**. It also achieves the highest average score across
 72 diverse reasoning benchmarks, improving over the base model by **12.1 points** and full privileged
 73 OPD by **2.3 points**.




Figure 2: AR-OPD constructs a dual-view anchored residual target. The **ANCHOR** supplies the reachable reference, while the **RESIDUAL** transfers controlled full-view guidance before distillation to the student.

2 Related Work

Knowledge distillation and on-policy distillation. Knowledge distillation trains a student to match a teacher distribution or teacher-generated outputs [Hinton et al., 2015, Kim and Rush, 2016, Sanh et al., 2019, Gu et al., 2023]. For autoregressive generation, off-policy distillation and SFT can suffer from exposure mismatch because training targets are not conditioned on student-generated prefixes [Ross et al., 2011, Lamb et al., 2016]. OPD addresses this by querying teachers on student-visited states and has become a strong post-training recipe for reasoning models [Thinking Machines Lab, 2025, Agarwal et al., 2024, Fu et al., 2026, Li et al., 2026]. Recent variants study objective stability, entropy control, and teacher-student compatibility, showing that dense teacher targets are useful but can also become harmful when teacher distributions are locally incompatible with the student [Jang et al., 2026, Jin et al., 2026, Yang et al., 2026b].

Privileged OPD. Privileged OPD conditions the teacher on auxiliary information such as demonstrations, reference traces, verifier feedback, or oracle answers [Lightman et al., 2023, Shenfeld et al., 2026, Zhao et al., 2026, Penaloza et al., 2026, Yang et al., 2025]. Existing variants differ mainly in how they stabilize this privileged teacher. SFT uses a delayed or EMA teacher to reduce training instability [Shenfeld et al., 2026], whereas OPSD keeps teacher and student synchronized and adds full-vocabulary JSD with per-token clipping [Zhao et al., 2026]. Despite these differences, they still treat the privileged distribution as an absolute imitation target, leaving target reachability and hindsight-induced support mismatch largely unaddressed.

3 Preliminaries

3.1 On-Policy Distillation

In autoregressive OPD, prompts $x \sim \mathcal{D}$ are paired with student rollouts $y = (y_1, \dots, y_T) \sim \pi_\theta(\cdot | x)$. The state at step t is $s_t = (x, y_{<t})$, and the student policy factorizes as

$$\pi_\theta(y | x) = \prod_{t=1}^T \pi_\theta(y_t | s_t). \quad (1)$$




Figure 3: Target-reliability diagnostics. **Left:** A token-level illustration of late-rollout teacher-student divergence, where disagreement concentrates near the final-answer region. **Middle:** Top- k disagreement rises again near the rollout tail and increases with privileged-context length. **Right:** No-overlap mass shows the same tail-end elevation, indicating that the teacher assigns probability mass to tokens outside the student’s local predictive support.

97 Given a teacher distribution $q(\cdot \mid s_t)$, the target of OPD is to minimize token-level KL divergence on
 98 student-visited states:

$$\mathcal{L}_{\text{OPD}}(\theta) = \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta} \left[\sum_{t=1}^T D_{\text{KL}}(q(\cdot \mid s_t) \parallel \pi_\theta(\cdot \mid s_t)) \right]. \quad (2)$$

99 Different OPD implementations can provide full logits, top- k logits, or hard next-token labels. We
 100 use the logit-level view because it makes target construction observable: changing the privileged
 101 context changes the entire teacher distribution, not only the sampled next token.

102 3.2 Privileged OPD

103 Privileged OPD conditions the teacher on auxiliary information z , such as a solution trace, an oracle
 104 answer, or an expert demonstration. With a frozen or delayed teacher parameter copy $\bar{\theta}$, the privileged
 105 teacher is

$$q_z(\cdot \mid s_t) = \pi_{\bar{\theta}}(\cdot \mid x, z, y_{<t}). \quad (3)$$

106 The student is trained without z at inference time. Direct privileged imitation therefore uses the
 107 training objective

$$\mathcal{L}_{\text{P-OPD}}(\theta) = \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_\theta} \left[\sum_{t=1}^T D_{\text{KL}}(q_z(\cdot \mid s_t) \parallel \pi_\theta(\cdot \mid s_t)) \right]. \quad (4)$$

108 4 Why Full-View Imitation Fails Local Alignment

109 Standard Privileged OPD assumes that oracle-augmented teachers provide uniformly reliable targets.
 110 We challenge this assumption by diagnosing the inherent reachability mismatch and signal fragility
 111 in absolute full-view imitation.

112 4.1 Target Unreachability: The Hindsight Support Gap

113 To quantify distributional misalignment, we measure Top- k Disagreement (D_k) and the Support Gap
 114 (M_τ^{out}). We compute these diagnostics on 1000 NuminaMath-style examples by sampling student
 115 rollouts from the public prompt and rescored the same generated prefixes under the student, partial-
 116 view teacher, and full-view teacher prompts. We report top- k disagreement with $k = 16$, and compute
 117 support gap as the teacher mass assigned outside the student’s locally plausible support; detailed
 118 estimator settings are provided in Section C.2. As shown in the middle and right panels of Figure 3,
 119 both metrics experience a tail-end elevation and scale with privileged-context length; the left panel
 120 illustrates the same late-rollout divergence pattern at the token level. Notably, for sequences exceeding
 121 1024 tokens, the support gap increases substantially in the final 40% of the rollout. This tail-end
 122 elevation is consistent with **hindsight-induced manifold drift**. One explanation is that in early
 123 reasoning steps, the student and teacher distributions often share common ground, but as generation
 124 deepens, the student’s prefix can accumulate slight deviations. The full-view teacher, conditioned on
 125 the gold future, remains anchored to the oracle trajectory. Rather than offering a localized correction,
 126 it assigns high probability to tokens that are justified by hindsight but logically disconnected from the

127 student’s actual accumulated prefix. Consequently, absolute full-view targets can pull the student
 128 toward **locally unsupported states**. Under the forward KL $D_{\text{KL}}(q \parallel \pi_\theta) = \sum_v q(v \mid s_t) \log \frac{q(v \mid s_t)}{\pi_\theta(v \mid s_t)}$,
 129 low-probability student terms $\pi_\theta(v \mid s_t)$ in the denominator can disproportionately dominate the
 130 expected loss. This can bias optimization toward oracle shortcuts instead of facilitating constructive,
 131 step-by-step reasoning transfer, consistent with recent analyses showing that privileged teacher-
 132 only signals can induce information leakage and that successful OPD depends on locally shared
 133 high-probability tokens at student-visited states [Yang et al., 2026a, Li et al., 2026].




Figure 4: Reliability of privileged teachers in long-horizon contexts. **Left:** Even with gold traces, the privileged teacher is imperfect and its reliability decays sharply as privileged context length increases. **Right:** The marginal gain of full-view over partial-view conditioning peaks at 256–512 tokens but then turns into negative marginal returns (up to -9.1%) in longer horizons. This highlights that full-view targets in long-horizon reasoning are often noisier than partial-view anchors.

134 4.2 Signal Fragility: The Illusion of Full-View Reliability

135 Gold privileged context improves average teacher behavior, but it does not make the full-
 136 view teacher a uniformly reliable token target. On 5,000 NuminaMath problems with
 137 Qwen2.5-7B-Instruct [Qwen Team, 2024], we run four rollouts for each teacher/student
 138 condition and extract final answers for scoring. Figure 4 shows that full-view teacher pass@4 drops
 139 from 90.18% in the < 128 -token bucket to 9.09% in the 768–1024-token bucket, and its marginal
 140 advantage over the partial view becomes negative by as much as -9.1% . This length sensitivity is
 141 consistent with recent OPD analyses showing that dense teacher targets can become locally incom-
 142 patible when teacher-side information exceeds the student’s reachable support [Li et al., 2026, Fu
 143 et al., 2026, Shenfeld et al., 2026, Penaloza et al., 2026].

144 The signal decomposition sharpens this point. Using the advantage scores defined in Equation (17),
 145 Figure A.1 shows that the partial privileged view is nearly as predictive of final correctness as the
 146 full-vs-student total advantage signal: partial advantage reaches ROC-AUC 0.766, close to the total
 147 signal at 0.779. By contrast, the marginal full-minus-partial signal falls to 0.585, indicating that
 148 the residual alone is a much weaker correctness classifier. Thus full-view information should not
 149 be treated as a standalone target; when used, it is better transferred as a controlled residual over a
 150 reliable partial anchor.

151 5 AR-OPD: Anchored Residual Guidance

152 AR-OPD replaces full-view imitation with a dual-view target. The method evaluates two teachers
 153 on the same student-generated state s_t : a partial teacher $q_{\text{part}}(\cdot \mid s_t)$ and a full teacher $q_{\text{full}}(\cdot \mid s_t)$.
 154 Figure 2 summarizes the target construction. The design follows a broader post-training pattern in
 155 which unstable long-horizon objects are decomposed into an anchor plus a controlled update, rather
 156 than optimized as a single monolithic target [Zhu et al., 2025, Li et al., 2025, Yang et al., 2026b].

157 5.1 Partial Privilege Defines the Local Anchor

158 Let the privileged information be $z = (z_1, \dots, z_L)$. We construct a partial privileged context by
 159 deterministic truncation:

$$z_{\text{part}} = z_{1:\lfloor \rho L \rfloor}, \quad \rho \in (0, 1). \quad (5)$$

160 The default setting is $\rho = 0.5$. The partial teacher is

$$q_{\text{part}}(\cdot | s_t) = \pi_{\bar{\theta}}(\cdot | x, z_{\text{part}}, y_{<t}). \quad (6)$$

161 The partial context exposes enough information to provide a structured setup while withholding later
 162 solution fragments, making the partial teacher a local anchor rather than an answer-leaking target.
 163 The front-partial choice is deliberately conservative: it resembles guided-prefix approaches that use
 164 early solution information to improve reachability, while avoiding direct supervision on the complete
 165 oracle trace [Qu et al., 2026, Shenfeld et al., 2026, Penaloza et al., 2026].

166 5.2 Full Privilege Enters as a Scaled Residual

167 The full teacher is

$$q_{\text{full}}(\cdot | s_t) = \pi_{\bar{\theta}}(\cdot | x, z, y_{<t}). \quad (7)$$

168 In log-probability space, the full-view learning signal decomposes as

$$\log q_{\text{full}} - \log \pi_{\theta} = (\log q_{\text{part}} - \log \pi_{\theta}) + (\log q_{\text{full}} - \log q_{\text{part}}). \quad (8)$$

169 The first term is the anchoring component, and the second term is the residual effect of giving the
 170 teacher the remaining privileged information.

Anchored residual target.

$$\log q_{\lambda}(v | s_t) = \underbrace{\log q_{\text{part}}(v | s_t)}_{\text{partial anchor}} + \underbrace{\lambda (\log q_{\text{full}}(v | s_t) - \log q_{\text{part}}(v | s_t))}_{\text{full-minus-partial residual}} - C_t. \quad (9)$$

171

Target construction at each student state s_t .

1. Query the **ANCHOR**, $q_{\text{part}}(\cdot | s_t)$, using the front partial privileged trace.
2. Query the **FULL VIEW**, $q_{\text{full}}(\cdot | s_t)$, using the complete privileged trace.
3. Transfer only $\lambda(\log q_{\text{full}} - \log q_{\text{part}})$, then normalize to obtain q_{λ} .
4. Distill the student against q_{λ} on the same on-policy prefix s_t .

172

173 The normalizer C_t makes $q_{\lambda}(\cdot | s_t)$ a valid distribution. The form makes the object-level change
 174 explicit: imitation no longer targets q_{full} directly; it targets an anchored belief update around q_{part} .
 175 The student is trained by KL distillation to the anchored residual target:

$$\mathcal{L}_{\text{AR-OPD}}(\theta) = \mathbb{E}_{x \sim \mathcal{D}, y \sim \pi_{\theta}} \left[\sum_{t=1}^T D_{\text{KL}}(q_{\lambda}(\cdot | s_t) \parallel \pi_{\theta}(\cdot | s_t)) \right]. \quad (10)$$

176 In implementation, targets are built on student-generated rollouts rather than fixed teacher traces.
 177 For each sampled prefix, the partial and full teacher prompts rescore the same next-token state with
 178 temperature 1.0; the anchored residual distribution is then formed in log space and optimized with
 179 forward KL. We use the same prompt and completion limits as the evaluation setup, and apply a
 180 small reference mixup and residual clipping only for numerical stability; the full configuration is
 181 listed in Table A.2.

182 The coefficient λ controls the amount of full-view residual transfer. When $\lambda = 0$, AR-OPD becomes
 183 partial privileged OPD; when $\lambda = 1$, it recovers the full privileged target by construction. Thus
 184 $\lambda < 1$ is a contractive residual regime that keeps the target anchored between the partial and full
 185 views, while $\lambda > 1$ extrapolates beyond the full-view log shift. Figure 5 visualizes this change:
 186 the full view is used as a directional residual around the partial anchor rather than as an absolute
 187 imitation target. The formulation makes λ closer to a residual-guidance strength than to a generic
 188 smoothing parameter, complementing work that studies reward extrapolation, target reformulation,
 189 and entropy-aware OPD control [Yang et al., 2026b, Jang et al., 2026, Jin et al., 2026].




Figure 5: Constructing reachable targets via anchored residual guidance. AR-OPD anchors on the student-compatible partial view and uses the full view only as a controlled directional residual.

6 Experiments: Performance Gains Follow Target Reliability

The evaluation separates outcome evidence from mechanism evidence through four checks: cross-domain outcome, anchor compatibility, residual value, and training dynamics.

6.1 Experimental Setup

Tasks and datasets. We evaluate four reasoning task groups. For mathematics, we train on NuminaMath-style CoT data [AI-MO/Numina, 2024] and evaluate on MATH500 [Hendrycks et al., 2021, Lightman et al., 2023] and AMC23 [Mathematical Association of America, 2023]. For code, we train on Magicoder-Evol-Instruct data [Wei et al., 2024, Luo et al., 2023] and evaluate HumanEval and MBPP with pass@1 [Chen et al., 2021, Austin et al., 2021]. For science, we use SciKnowEval Chemistry L-3 [Feng et al., 2024] with exact-match multiple-choice accuracy. For medical QA, we train on MedMCQA [Pal et al., 2022] and evaluate on MedQA [Jin et al., 2020] and MedCQA [Pal et al., 2022]. We use the language-model evaluation harness for standardized evaluation plumbing where applicable [Gao et al., 2024].

Model and training. Unless stated otherwise, we use Qwen2.5-7B-Instruct as the base model [Qwen Team, 2024]. The student and teacher views share the same model family and differ only in the privileged context exposed to the teacher during target construction. Across tasks, we set the maximum completion length to 1024 tokens and use a 2048-token context window including the input prompt. The default partial rate is $\rho = 0.5$. We sweep $\lambda \in \{0.4, 0.6, 0.8, 1.0, 1.2\}$ to study residual strength, keeping other hyperparameters fixed across methods.

Compared methods. We compare the base model, supervised fine-tuning (SFT), full privileged OPD, partial privileged OPD, and AR-OPD. The comparison separates demonstration-only adaptation, full privileged self-distillation, and guided partial-privilege training [Ouyang et al., 2022, Shenfeld et al., 2026, Qu et al., 2026]. Full OPD trains directly against q_{full} with an EMA self-teacher. AR-OPD trains against q_λ , the anchor-residual target.

Evaluation protocol. All reported methods are evaluated with the same public prompts and without privileged context at inference time. Privileged information is used only to construct training targets, so the comparison isolates how different target constructions affect the same base model. For math and multiple-choice tasks, we extract the final answer with the task-specific evaluator; for code, we use the standard unit-test pass@1 protocol. All distillation variants use the same on-policy rollout budget, optimizer settings, prompt limits, and final-checkpoint reporting rule. Thus, differences in the main table reflect how the target distribution is constructed rather than changes in data scale, inference-time context, or evaluator choice.

6.2 Main Results: Performance

Table 1: Anchored residual guidance achieves the strongest average performance across domains. Metrics are accuracy or pass@1; code, science, and medical results remain part of the main cross-domain evidence, while checkpointed math-only diagnostics are separated in Section D.

Domain	Benchmark	Base	SFT	Partial OPD	Full OPD	AR-OPD
Math	MATH500	61.8	63.2	70.6	71.0	74.6
	AMC23	45.0	47.5	45.0	55.0	57.5
Code	HumanEval	79.2	76.8	79.4	78.8	82.6
	MBPP	77.5	79.4	77.8	78.5	80.4
Science	SciKnowEval	32.12	47.4	64.2	68.0	68.5
Medical	MedQA	58.7	66.0	64.3	65.2	66.7
	MedCQA	53.2	56.7	55.4	59.0	61.5
Average	All benchmarks	58.2	62.4	65.2	67.9	70.3

Performance. AR-OPD consistently improves over standard SFT and competitive on-policy distillation baselines across reasoning-intensive and knowledge-heavy domains. As shown in Table 1, AR-OPD obtains the best average score of 70.3, improving over the base model by **12.1 points**, SFT by **7.9 points**, and the strong Full OPD baseline by **2.3 points**. The gains are largest in mathematical reasoning, where hindsight-conditioned targets are most likely to induce support mismatch: AR-OPD reaches 74.6 on MATH500 and 57.5 on AMC23, exceeding Full OPD by **3.6 and 2.5 points**, respectively. Notably, even Partial OPD, which remains blinded to the final answer, outperforms standard full-trace SFT by an average of 2.8 points. While Partial OPD lacks destination awareness, AR-OPD adds residual guidance over a locally valid partial anchor, yielding a **5.1-point average gain over Partial OPD**. The improvement extends beyond math: AR-OPD also ranks first on HumanEval, MBPP, SciKnowEval, MedQA, and MedCQA.




Figure 6: AR-OPD improves validation accuracy, reduces shortcuts, and is strongest on long rollouts. Left: Validation accuracy during Math10K training: AR-OPD shows consistent learning gains, while Full OPD peaks early and then largely plateaus. Middle: Shortcut-generation count over training steps; at the final checkpoint / percentile (100%), AR-OPD has 18 shortcut events versus 23 for Full OPD, a 21.7% relative reduction. Right: Accuracy grouped by rollout length, where AR-OPD improves over Full OPD by 5.1 points on 512–768-token rollouts and 7.2 points on 768–1024-token rollouts.

6.3 Shortcut Dynamics and Long-Rollout Accuracy

Shortcut dynamics. The shortcut diagnostic quantifies the failure mode from Figure 1. By anchoring to the partial view, AR-OPD consistently suppresses shortcut events, reducing the average count across all 21 training checkpoints from 12.95 for Full OPD to 7.52. At the final checkpoint, AR-OPD produces only 18 shortcuts compared to Full OPD’s 23, a **21.7% relative reduction**.

Long-horizon accuracy. AR-OPD’s advantage scales with trajectory length (Figure 6, right). While competitive on short rollouts, it surpasses Full OPD by **5.1 points** on 512–768-token rollouts and **7.2 points** on 768–1024-token rollouts. These results suggest that, as reasoning horizons grow, anchored residual targets reduce the compounding hindsight drift that can degrade monolithic full-view imitation.

Empirical highlight. The results show that controlling the full-view residual improves over full privileged OPD, while the gain over Partial OPD confirms that the residual itself adds useful destination-directed signal.

6.4 Lambda Sweep and Residual Control

Table 2: Lambda sweep for the AR-OPD residual scale on the Math-10K checkpointed run. Metrics are relaxed accuracy. Contractive residual transfer ($\lambda < 1$) gives the strongest settings, while extrapolating past the full-view residual with $\lambda = 1.2$ degrades performance.

Residual scale	Regime	Math500	AMC23
$\lambda = 0.4$	Weak residual	71.2	50.0
$\lambda = 0.6$	Contractive residual	71.6	57.5
$\lambda = 0.8$	Contractive residual	74.6	55.0
$\lambda = 1.0$	Full-view limit	71.8	55.0
$\lambda = 1.2$	Extrapolated residual	69.0	50.0

Residual scale determines how aggressively AR-OPD transfers the full-view log shift. By Equation (9), $\lambda = 1$ recovers q_{full} , values below one keep the target contractively tied to the partial anchor, and values above one extrapolate beyond the full-view direction. The sweep in Table 2 shows that $\lambda = 0.8$ is strongest on Math500 and $\lambda = 0.6$ is strongest on AMC23, whereas the weak residual setting $\lambda = 0.4$ underuses answer-directed guidance and the extrapolated setting $\lambda = 1.2$ underperforms both contractive settings. We therefore treat λ as the main control knob balancing partial-anchor stability against full-view over-transfer, in line with recent OPD and RL-for-reasoning work showing that target scale and reformulation materially affect optimization behavior [Yang et al., 2026b, Jin et al., 2026, Wang et al., 2025].

7 Discussion: Decomposing Privileged Distribution Shift

Our results frame privilege as a target-design problem rather than a binary choice between using or discarding oracle information. Full-view oracle distributions can provide useful destination-directed signal, but absolute imitation can create a hindsight trap: tokens valid under the privileged future may remain unsupported from the student’s current prefix. The diagnostics suggest that this mismatch grows near rollout tails and that longer privileged context does not necessarily improve teacher reliability. AR-OPD mitigates this by separating supervision into a locally valid anchor and a controlled residual, turning full-view information into a bounded update rather than a standalone coordinate.

Future work. Promising directions include adaptive anchors based on student uncertainty or gradient variance, broader privileged modalities beyond gold traces such as search heuristics or human preferences, and long-horizon settings such as multi-turn agents or repository-level coding where hindsight interference may be more pronounced.

8 Limitations

AR-OPD requires two teacher forward passes per student state, increasing target-generation cost relative to single-view OPD and motivating comparison with search- or verifier-based alternatives [Liu et al., 2025, Yang et al., 2025, Walder and Karkhanis, 2025]. Our experiments use one open model family, Qwen2.5-7B-Instruct, single seeds, and modest data sizes; broader sweeps over data, scale, and families such as Llama [Touvron et al., 2023] remain important. Algorithmically, we use a fixed partial anchor ($\rho = 0.5$) and static residual scale; adaptive schedules, diagnostics beyond NuminaMath, and structured-output robustness remain open [Liu et al., 2023, Hassid et al., 2025, Madaan et al., 2023, Yuan et al., 2024].

References

- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos Garea, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *The Twelfth International Conference on Learning Representations*, 2024.
- AI-MO/Numina. NuminaMath. <https://huggingface.co/datasets/AI-MO/NuminaMath-CoT>, 2024.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Arwen Bradley and Preetum Nakkiran. Classifier-free guidance is a predictor-corrector, 2024. URL <https://arxiv.org/abs/2408.09000>.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde De Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Hanze Dong, Wei Xiong, Bo Pang, Haoxiang Wang, Han Zhao, Yingbo Zhou, Nan Jiang, Doyen Sahoo, Caiming Xiong, and Tong Zhang. Rlhf workflow: From reward modeling to online rlhf. *arXiv preprint arXiv:2405.07863*, 2024.
- Kehua Feng, Keyan Ding, Weijie Wang, Xiang Zhuang, Zeyuan Wang, Ming Qin, Yu Zhao, Jianhua Yao, Qiang Zhang, and Huajun Chen. Sciknoweval: Evaluating multi-level scientific knowledge of large language models. *arXiv preprint arXiv:2406.09098*, 2024.
- Yuqian Fu, Haohuan Huang, Kaiwen Jiang, Jiakai Liu, Zhuo Jiang, Yuanheng Zhu, and Dongbin Zhao. Revisiting on-policy distillation: Empirical failure modes and simple fixes. *arXiv preprint arXiv:2603.25562*, 2026.
- Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, Sid Black, Anthony DiPofi, Charles Foster, Laurence Golding, Jeffrey Hsu, Alain Le Noac’h, Haonan Li, Kyle McDonell, Niklas Muennighoff, Chris Ociepa, Jason Phang, Laria Reynolds, Hailey Schoelkopf, Aviya Skowron, Lintang Sutawika, Eric Tang, Anish Thite, Ben Wang, Kevin Wang, and Andy Zou. The language model evaluation harness, 07 2024. URL <https://zenodo.org/records/12608602>.
- GLM-5 Team. Glm-5: from vibe coding to agentic engineering, 2026. URL <https://arxiv.org/abs/2602.15763>.
- Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. Minillm: Knowledge distillation of large language models. *arXiv preprint arXiv:2306.08543*, 2023.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Michael Hassid, Gabriel Synnaeve, Yossi Adi, and Roy Schwartz. Don’t overthink it. preferring shorter thinking chains for improved llm reasoning. *arXiv preprint arXiv:2505.17813*, 2025.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2021. URL <https://openreview.net/forum?id=7Bywt2mQsCe>.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. In *NeurIPS 2022 Workshop on Score-Based Methods*, 2022.

325 Ijun Jang, Jewon Yeom, Juan Yeo, Hyunggu Lim, and Taesup Kim. Stable on-policy distillation
326 through adaptive target reformulation. *arXiv preprint arXiv:2601.07155*, 2026.

327 Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. What
328 disease does this patient have? a large-scale open domain question answering dataset from medical
329 exams. *arXiv preprint arXiv:2009.13081*, 2020.

330 Woogyeol Jin, Taywon Min, Yongjin Yang, Swanand Ravindra Kadhe, Yi Zhou, Dennis Wei, Nathalie
331 Baracaldo, and Kimin Lee. Entropy-aware on-policy distillation of language models. *arXiv preprint
332 arXiv:2603.07079*, 2026.

333 Yoon Kim and Alexander M Rush. Sequence-level knowledge distillation. In *Proceedings of the
334 2016 conference on empirical methods in natural language processing*, pages 1317–1327, 2016.

335 Alex M Lamb, Anirudh Goyal ALIAS PARTH GOYAL, Ying Zhang, Saizheng Zhang, Aaron C
336 Courville, and Yoshua Bengio. Professor forcing: A new algorithm for training recurrent networks.
337 *Advances in neural information processing systems*, 29, 2016.

338 Yaxuan Li, Yuxin Zuo, Bingxiang He, Jinqian Zhang, Chaojun Xiao, Cheng Qian, Tianyu Yu, Huan-
339 ang Gao, Wenkai Yang, Zhiyuan Liu, and Ning Ding. Rethinking on-policy distillation of large
340 language models: Phenomenology, mechanism, and recipe. *arXiv preprint arXiv:2604.13016*,
341 2026.

342 Zhuofeng Li, Haoxiang Zhang, Seungju Han, Sheng Liu, Jianwen Xie, Yu Zhang, Yejin Choi, James
343 Zou, and Pan Lu. In-the-flow agentic system optimization for effective planning and tool use.
344 *arXiv preprint arXiv:2510.05592*, 2025.

345 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
346 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint
347 arXiv:2305.20050*, 2023.

348 Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated
349 by chatGPT really correct? rigorous evaluation of large language models for code generation.
350 In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL <https://openreview.net/forum?id=lqvxx610Cu7>.
351

352 Zijun Liu, Peiyi Wang, Runxin Xu, Shirong Ma, Chong Ruan, Peng Li, Yang Liu, and Yu Wu.
353 Inference-time scaling for generalist reward modeling. *arXiv preprint arXiv:2504.02495*, 2025.

354 Ziyang Luo, Can Xu, Pu Zhao, Qingfeng Sun, Xiubo Geng, Wenxiang Hu, Chongyang Tao, Jing
355 Ma, Qingwei Lin, and Daxin Jiang. Wizardcoder: Empowering code large language models with
356 evol-instruct, 2023. URL <https://arxiv.org/abs/2306.08568>.

357 Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri
358 Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement
359 with self-feedback. *arXiv preprint arXiv:2303.17651*, 2023.

360 Mathematical Association of America. 2023 american mathematics com-
361 petitions problems. [https://maa.org/math-competitions/
362 american-mathematics-competitions-amc/](https://maa.org/math-competitions/american-mathematics-competitions-amc/), 2023.

363 Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong
364 Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow
365 instructions with human feedback. *Advances in Neural Information Processing Systems*, 2022.

366 Ankit Pal, Logesh Kumar Umapathi, and Malaikannan Sankarasubbu. Medmcqa: A large-scale
367 multi-subject multi-choice dataset for medical domain question answering. In *Proceedings of the
368 Conference on Health, Inference, and Learning*, volume 174 of *Proceedings of Machine Learning
369 Research*, pages 248–260. PMLR, 2022. URL [https://proceedings.mlr.press/
370 v174/pal22a.html](https://proceedings.mlr.press/v174/pal22a.html).

371 Emiliano Penalosa, Dheeraj Vattikonda, Nicolas Gontier, Alexandre Lacoste, Laurent Charlin,
372 and Massimo Caccia. Privileged information distillation for language models. *arXiv preprint
373 arXiv:2602.04942*, 2026.

374 Yuxiao Qu, Amrith Setlur, Virginia Smith, Ruslan Salakhutdinov, and Aviral Kumar. Pope: Learning
375 to reason on hard problems via privileged on-policy exploration. *arXiv preprint arXiv:2601.18779*,
376 2026.

377 Qwen Team. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

378 Qwen Team. Qwen3 technical report, 2025. URL <https://arxiv.org/abs/2505.09388>.

379 Stephane Ross, Geoffrey J. Gordon, and J. Andrew Bagnell. A reduction of imitation learning
380 and structured prediction to no-regret online learning. In *International Conference on Artificial*
381 *Intelligence and Statistics*, 2011.

382 Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of
383 bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.

384 John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy
385 optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.

386 Idan Shenfeld, Mehul Damani, Jonas Hübner, and Pulkit Agrawal. Self-distillation enables continual
387 learning. *arXiv preprint arXiv:2601.19897*, 2026.

388 Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng,
389 Haibin Lin, and Chuan Wu. Hybridflow: A flexible and efficient rlhf framework. *arXiv preprint*
390 *arXiv: 2409.19256*, 2024.

391 Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy
392 Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following llama model.
393 <https://crfm.stanford.edu/2023/03/13/alpaca.html>, 2023.

394 Thinking Machines Lab. On-policy distillation. [https://thinkingmachines.ai/blog/](https://thinkingmachines.ai/blog/on-policy-distillation/)
395 [on-policy-distillation/](https://thinkingmachines.ai/blog/on-policy-distillation/), 2025.

396 Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothee
397 Lacroix, Baptiste Roziere, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand
398 Joulin, Edouard Grave, and Guillaume Lample. Llama: Open and efficient foundation language
399 models. *arXiv preprint arXiv:2302.13971*, 2023.

400 Christian Walder and Deep Karkhanis. Pass@ k policy optimization: Solving harder reinforcement
401 learning problems. *arXiv preprint arXiv:2505.15201*, 2025.

402 Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen,
403 Jianxin Yang, Zhenru Zhang, et al. Beyond the 80/20 rule: High-entropy minority tokens drive
404 effective reinforcement learning for llm reasoning. *arXiv preprint arXiv:2506.01939*, 2025.

405 Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny
406 Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in*
407 *Neural Information Processing Systems*, 35:24824–24837, 2022.

408 Yuxiang Wei, Zhe Wang, Jiawei Liu, Yifeng Ding, and Lingming Zhang. Magicoder: Empowering
409 code generation with OSS-instruct. In *Proceedings of the 41st International Conference on Machine*
410 *Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 52632–52657. PMLR,
411 21–27 Jul 2024. URL <https://proceedings.mlr.press/v235/wei24h.html>.

412 Xiaomi MiMo Team. Mimo: Unlocking the reasoning potential of language model – from pretraining
413 to posttraining, 2025. URL <https://arxiv.org/abs/2505.07608>.

414 Chenxu Yang, Chuanyu Qin, Qingyi Si, Minghui Chen, Naibin Gu, Dingyu Yao, Zheng Lin, Weiping
415 Wang, Jiaqi Wang, and Nan Duan. Self-distilled rlvr. *arXiv preprint arXiv:2604.03128*, 2026a.

416 Wenkai Yang, Jingwen Chen, Yankai Lin, and Ji-Rong Wen. Deepcritic: Deliberate critique with
417 large language models. *arXiv preprint arXiv:2505.00662*, 2025.

418 Wenkai Yang, Weijie Liu, Ruobing Xie, Kai Yang, Saiyong Yang, and Yankai Lin. Learning
419 beyond teacher: Generalized on-policy distillation with reward extrapolation. *arXiv preprint*
420 *arXiv:2602.12125*, 2026b.

- 421 Weizhe Yuan, Richard Yuanzhe Pang, Kyunghyun Cho, Sainbayar Sukhbaatar, Jing Xu, and Jason
422 Weston. Self-rewarding language models. *arXiv preprint arXiv:2401.10020*, 2024.
- 423 Siyan Zhao, Zhihui Xie, Mengchen Liu, Jing Huang, Guan Pang, Feiyu Chen, and Aditya Grover.
424 Self-distilled reasoner: On-policy self-distillation for large language models. *arXiv preprint*
425 *arXiv:2601.18734*, 2026.
- 426 He Zhu, Junyou Su, Peng Lai, Ren Ma, Wenjia Zhang, Linyi Yang, and Guanhua Chen. Anchored
427 supervised fine-tuning. *arXiv preprint arXiv:2509.23753*, 2025.

428 Use of LLMs

429 Large language models were utilized as auxiliary tools to assist with code writing, manuscript editing,
430 and figure preparation. The authors maintained full review and control over all research ideas,
431 experimental designs, analyses, and final content. LLM assistance was used for LaTeX cleanup,
432 wording refinement, caption compression, and drafting small analysis scripts during revision. LLMs
433 were not used to generate benchmark measurements; all reported numbers come from the described
434 experiments and were checked by the authors.

435 A Prompt for Privileged Information

436 For each training example, we construct a standard student prompt and two privileged teacher
437 prompts. The student prompt contains only the original programming problem. The privileged
438 prompts additionally include an example response before asking the model to produce its own
439 response.

Template for full and partial privileged information.

```
[Original problem]
This is an example for a response to the question:
[Privileged response]
Now answer with a response of your own, including the thinking
process.
```

440

441 For the full privileged condition, [Privileged response] is the complete reference trace,
442 including the final code answer when available. For the partial privileged condition, [Privileged
443 response] is a partial trace that removes the final answer. In the code data, the partial trace is
444 constructed by removing fenced code blocks and inline code snippets from the reference response,
445 then truncating the remaining explanation to the first 50% of characters without splitting a word.

Resulting training record.

- prompt: the original programming problem;
- partial_teacher_prompt: the original problem plus partial privileged information;
- full_teacher_prompt: the original problem plus full privileged information;
- gold_trace: the full reference response;
- partial_trace: the partial privileged response;
- gold_answer: the extracted final code answer.

446

447 B Training and Evaluation Provenance

448 To make the main comparison reproducible, Tables A.1 to A.3 summarize the training data, shared
449 optimization configuration, and evaluation protocol used for Table 1. Unless otherwise specified, all
450 methods use seed 42 and report the final checkpoint.

Table A.1: Training data and privileged-view construction. Each row summarizes the source and dual-view construction used to build the corresponding training records.

Domain	Train size	Source	View construction
Math	10K	NuminaMath-style CoT subset	Partial teacher keeps the first 50% of the reasoning trace at a complete-word boundary and removes the final answer; full teacher keeps the complete trace with the gold final answer.
Code	4K	Magocoder-Evol-Instruct-110K [Wei et al., 2024]; upstream Evol-CodeAlpaca-style data	Partial trace removes fenced code blocks and in-line code snippets, then keeps the first 50% of the remaining explanation without splitting a word; full trace keeps the complete reference response.
Science	2K	SciKnowEval Chemistry L-3 [Feng et al., 2024]	Multiple-choice demonstrations use the same dual-view construction; accuracy is computed by exact match on the final answer choice.
Medical	2K	MedMCQA subset [Pal et al., 2022]	Multiple-choice demonstrations use the same dual-view construction; accuracy is computed by exact match on the final answer choice.

Table A.2: Shared dual-GPU serving and training configuration. Code uses 4K training examples; science and medical use 2K examples.

Parameter	Value
Base model	Qwen2.5-7B-Instruct
Hardware topology	GPU 0 training, GPU 1 vLLM teacher serving
Learning rate / scheduler	2×10^{-5} , cosine, warmup ratio 0.1
Epochs / max steps	1 epoch, epoch-based training
Per-device batch / global batch	1 / 64
Gradient accumulation	64
Prompt / completion length	max prompt 1024, max completion 1024
Context window	vLLM max model length 2048, including input prompt
Training temperature	1.0
Precision	bf16 when GPU is available
Seed	42
Distillation objective	forward KL, distill alpha 0.0
Ref mixup / residual clip	mixup alpha 0.01, residual clip 5.0
Default residual scale	$\lambda = 0.6$, with sweep in Table 2
Teacher CPU offload	enabled via dual-teacher CPU offload

Table A.3: Evaluation protocol for the main results. All methods report final-checkpoint performance with seed 42.

Benchmark	Decoding	Evaluator	Metric
MATH500	Greedy, temperature 0.0	Math answer parser with relaxed normalization	Relaxed accuracy
AMC23	Greedy, temperature 0.0	Math answer parser with relaxed normalization	Relaxed accuracy
HumanEval	Greedy, temperature 0.0	Original HumanEval evaluator [Chen et al., 2021]	pass@1
MBPP	Greedy, temperature 0.0	Original MBPP evaluator [Austin et al., 2021]	pass@1
SciKnowEval	Greedy, temperature 0.0	Exact match on final multiple-choice answer	Accuracy
MedQA	Greedy, temperature 0.0	Exact match on final multiple-choice answer	Accuracy
MedCQA	Greedy, temperature 0.0	Exact match on final multiple-choice answer	Accuracy

C Diagnostic Definitions

C.1 Connection to Classifier-Free Guidance

Classifier-Free Guidance (CFG) combines unconditional and conditional score estimates by scaling the residual effect of a condition [Ho and Salimans, 2022]. In diffusion models, this can be written abstractly as

$$\tilde{s}_\gamma(x_t, c) = s_\theta(x_t, \emptyset) + \gamma(s_\theta(x_t, c) - s_\theta(x_t, \emptyset)), \quad (11)$$

where the unconditional score provides a reference trajectory and the conditional residual steers sampling. Recent analyses further connect CFG to predictor-corrector behavior [Bradley and Nakkiran, 2024]. We use this only as a loose analogy: AR-OPD does not import a diffusion objective, but it adopts the same structural separation between an anchor distribution and a scaled directional residual.

C.2 Diagnostic Estimator

We estimate the target-reliability diagnostics on 1000 NuminaMath-style math examples. For each example, we sample a student rollout from the public prompt and treat every generated prefix as a student-visited state s_t . We then rescore the same generated token sequence under the public student prompt, partial privileged teacher prompt, and full privileged teacher prompt using the same checkpoint. Logits are converted to probabilities with temperature $T = 1.0$; generation uses top- $p = 1.0$, maximum generation length 1024, and maximum prompt length 2048. We report top- k disagreement with $k = 16$, averaged over token positions and grouped by five rollout-percentile bins and privileged-context-length buckets.

For a teacher distribution q , we first define the top- k overlap ratio as

$$O_k(q, \pi_\theta; s_t) = \frac{|\text{Top}_k(q(\cdot | s_t)) \cap \text{Top}_k(\pi_\theta(\cdot | s_t))|}{k}. \quad (12)$$

The corresponding top- k disagreement ratio is

$$D_k(q, \pi_\theta; s_t) = 1 - O_k(q, \pi_\theta; s_t). \quad (13)$$

A larger D_k indicates weaker local alignment between the teacher’s high-probability tokens and the student’s reachable predictive support on the same prefix.

The implementation-level diagnostic also records shared teacher mass over the overlapping top- k set:

$$\hat{B}_k(q, \pi_\theta; s_t) = \sum_{v \in \text{Top}_k(q(\cdot | s_t)) \cap \text{Top}_k(\pi_\theta(\cdot | s_t))} q(v | s_t). \quad (14)$$

The corresponding top- k no-overlap mass proxy is the teacher mass assigned to top- k teacher tokens that do not appear in the student’s top- k set:

$$\hat{N}_k(q, \pi_\theta; s_t) = \sum_{v \in \text{Top}_k(q(\cdot | s_t)) \setminus \text{Top}_k(\pi_\theta(\cdot | s_t))} q(v | s_t). \quad (15)$$

This top- k proxy should be distinguished from the full-vocabulary thresholded support gap below.

We also define the teacher mass outside the student’s local predictive support at threshold τ :

$$M_\tau^{\text{out}}(q, \pi_\theta; s_t) = \sum_{v: \pi_\theta(v | s_t) < \tau} q(v | s_t). \quad (16)$$

This no-overlap mass, also referred to as the support gap, measures how much probability the teacher assigns to tokens that the student currently treats as locally implausible. Equivalently, since q is normalized, $M_\tau^{\text{out}} = 1 - \sum_{v: \pi_\theta(v | s_t) \geq \tau} q(v | s_t)$. Thus the plotted no-overlap quantity is already the complement of in-support teacher mass. High values indicate that the teacher distribution may be correct under its privileged view but locally difficult for the student to imitate from s_t .

483 C.3 Teacher-Signal Advantage Scores

484 To measure which part of the privileged signal predicts final-answer correctness, we compute log-
 485 probability advantage scores over a token span $\mathcal{T}$. The total, partial, and marginal advantages
 486 are

$$\begin{aligned} A_{\text{full}}(\mathcal{T}) &= \sum_{t \in \mathcal{T}} (\log q_{\text{full}}(y_t | s_t) - \log \pi_{\theta}(y_t | s_t)), \\ A_{\text{part}}(\mathcal{T}) &= \sum_{t \in \mathcal{T}} (\log q_{\text{part}}(y_t | s_t) - \log \pi_{\theta}(y_t | s_t)), \\ A_{\text{marg}}(\mathcal{T}) &= A_{\text{full}}(\mathcal{T}) - A_{\text{part}}(\mathcal{T}) = \sum_{t \in \mathcal{T}} (\log q_{\text{full}}(y_t | s_t) - \log q_{\text{part}}(y_t | s_t)). \end{aligned} \quad (17)$$

487 We use each scalar score as a predictor of final-answer correctness and report ROC-AUC in Figure A.1.

488 C.4 Shortcut-Event Counting Protocol

489 To quantify privileged-information leakage in Figure 6, we evaluate trained privileged-OPD check-
 490 points on a held-out NuminaMath validation split and generate 5K responses per checkpoint. We use
 491 a deterministic response-level detector for shortcut-like evidence, such as references to unavailable
 492 solutions, answer-conditioned certainty without local derivation, or prompt-external answer cues.
 493 A response is counted once if any rule fires, so the reported count is a conservative response-level
 494 diagnostic of target-side leakage rather than a judge of mathematical correctness.

495 D Math Deep Dive and Training Diagnostics

496 Section D reports supplemental diagnostics for the math setting; these results do not replace the main
 497 cross-domain comparison in Table 1. The Numina500 split exposes checkpoint and residual-scale
 498 sensitivity rather than redefining the headline benchmark table.

499 D.1 Teacher-Signal Diagnostics

500 Figure A.1 and Figure A.2 provide the appendix-
 501 level diagnostics behind the main mechanism
 502 claim: the partial view carries stable correctness-
 503 predictive signal, while the full-view target
 504 creates larger disagreement and support gaps
 505 against the student.

506 D.2 Checkpoint-Level Math Results

507 The checkpoint sweep is used only as supple-
 508 mental evidence for residual-scale sensitivity:
 509 contractive AR-OPD settings are strongest on
 510 the headline math diagnostics, while the main
 511 cross-domain comparison remains the decisive
 512 result.

513 D.3 Failure Modes

514 Some raw evaluation failures are amplified by
 515 output-length inflation and answer-format insta-
 516 bility rather than pure reasoning failure. Relaxed
 517 grading recovers many errors, especially for the
 518 partial-only model, suggesting that structured-
 519 answer compliance remains a bottleneck and
 520 that relaxed gains should be interpreted as for-
 521 matting sensitivity rather than genuine reasoning
 522 improvement.




Figure A.1: Correctness-predictive teacher signal. The partial privileged signal nearly matches the total signal in ROC-AUC, while the marginal full-minus-partial residual is substantially weaker.




Figure A.2: Teacher-student support-gap diagnostics. Full-view targets show higher top- k disagreement and no-overlap mass than partial-view anchors across rollout positions.